

Distributed Consistency with CRDTs

- Kishan Sagathiya, @kishansagathiya

Software Engineer at Protocol Labs,
Member of IPFS


Conflict free

Replicated

Data-Types

Modified by many, but eventually
consistent

Distributed Databases



Modified by many, but eventually
consistent

Collaborative Text Editors

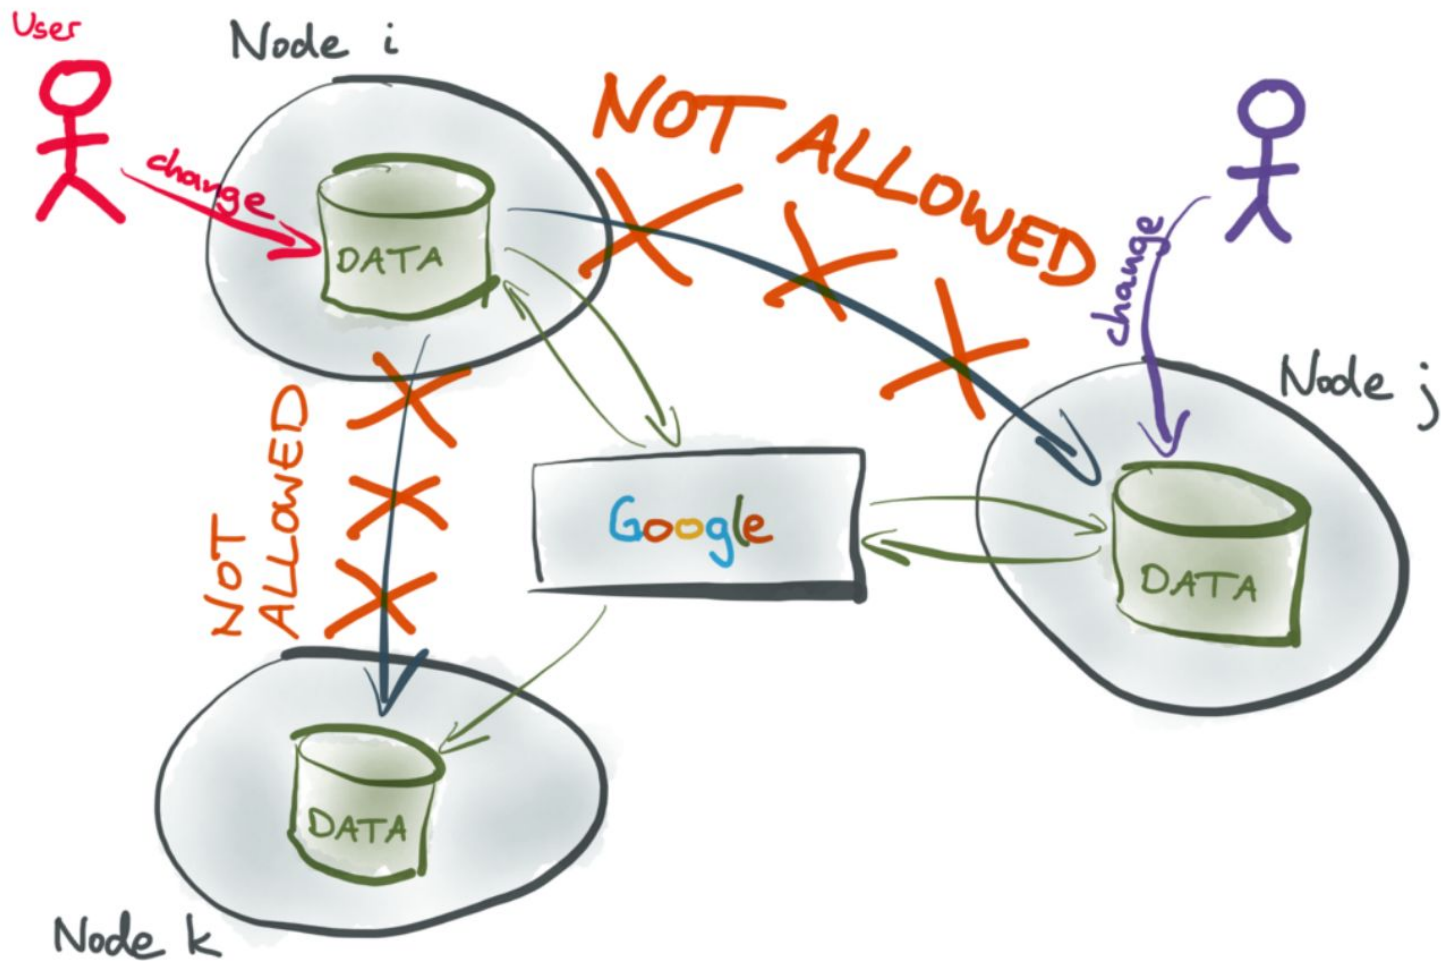


OPERATIONAL TRANSFORMATION (OT)

- e.g. Google Docs,

MS Office Online

OPERATIONAL TRANSFORMATION IN GOOGLE DOCS.



- So, *Operational Transformation* requires a server
- Can we do it *without a server*? YES

<https://peerpad.net/>

CRDTs



PeerPad^α

What have people built?

- *Redis*: distributed, highly available and scalable in-memory database
- *Automerger*: A JSON-like data structure (a CRDT) that can be modified concurrently by different users, and merged again automatically.
- *Orbitdb*: Peer-to-Peer Databases for the Decentralized Web
- *Riak*: decentralized datastore
- *PeerPad*: is a real-time collaborative text editor
- *TomTom GPS* uses it for data synchronization
- *Teletype for atom*: collaborate on code in real time
- *Chat in League of Legends*
- *Cosmos DB* by Microsoft

And other things....

CRDTs are data types which provide *strong eventual consistency* among different *replicas* in a distributed system by requiring some properties from the *state* and/or the *operations* applied to modify it.

Strong Eventual Consistency

If two replicas have received the *same updates*,
their state will be the *same*

State based CRDTs (*Convergent* CRDTs)

Operation based CRDTs (*Commutative* CRDTs)

Operation based CRDTs

Operations that modify states must be *commutative*

$$\begin{array}{l} A - 3 \\ + 4 \end{array}$$

$$\begin{array}{l} A + 4 \\ - 3 \end{array}$$

✓

$$\begin{array}{l} A - 3 \\ * 4 \end{array}$$

$$\begin{array}{l} A * 4 \\ - 3 \end{array}$$

✗

Operation based CRDTs

Exactly once delivery semantics

Sum(A, 3)

X

Max(A, 3)

✓

State based CRDTs

In state-based CRDTs, the states in different replicas and different moments form a *monotonic join semilattice*.



JOIN SEMILATTICE YOU SAY!



TELL ME MORE ABOUT THAT

- less than or equal to

$a \leq b$ or $b \leq a$

- incomparable

$a \parallel b$

- join

$a \vee b$


- An *order* is a binary relation \leq on a set S, written $\langle S, \leq \rangle$
- examples

less than or equal to $2 \leq 4$


descendent-of daughter \leq mother

- *Total Order*

Comes-before order



Less Than or Equal To
Order

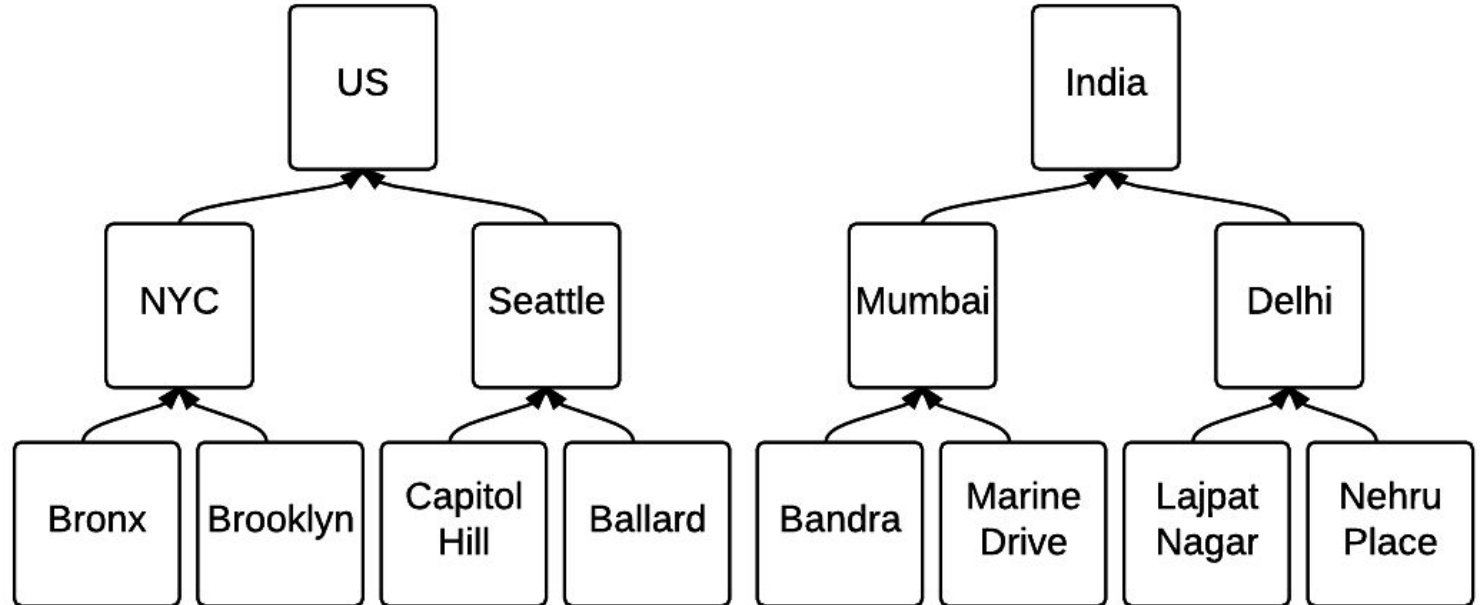


- *Partial Order*

Seattle \leq US and Brooklyn \leq US


Seattle // NYC and Bronx // Mumbai

Located-in Order





- A *vector clock timestamp* is a collection of logical timestamps for all the nodes or processes we're interested in.

Happened-Before Order



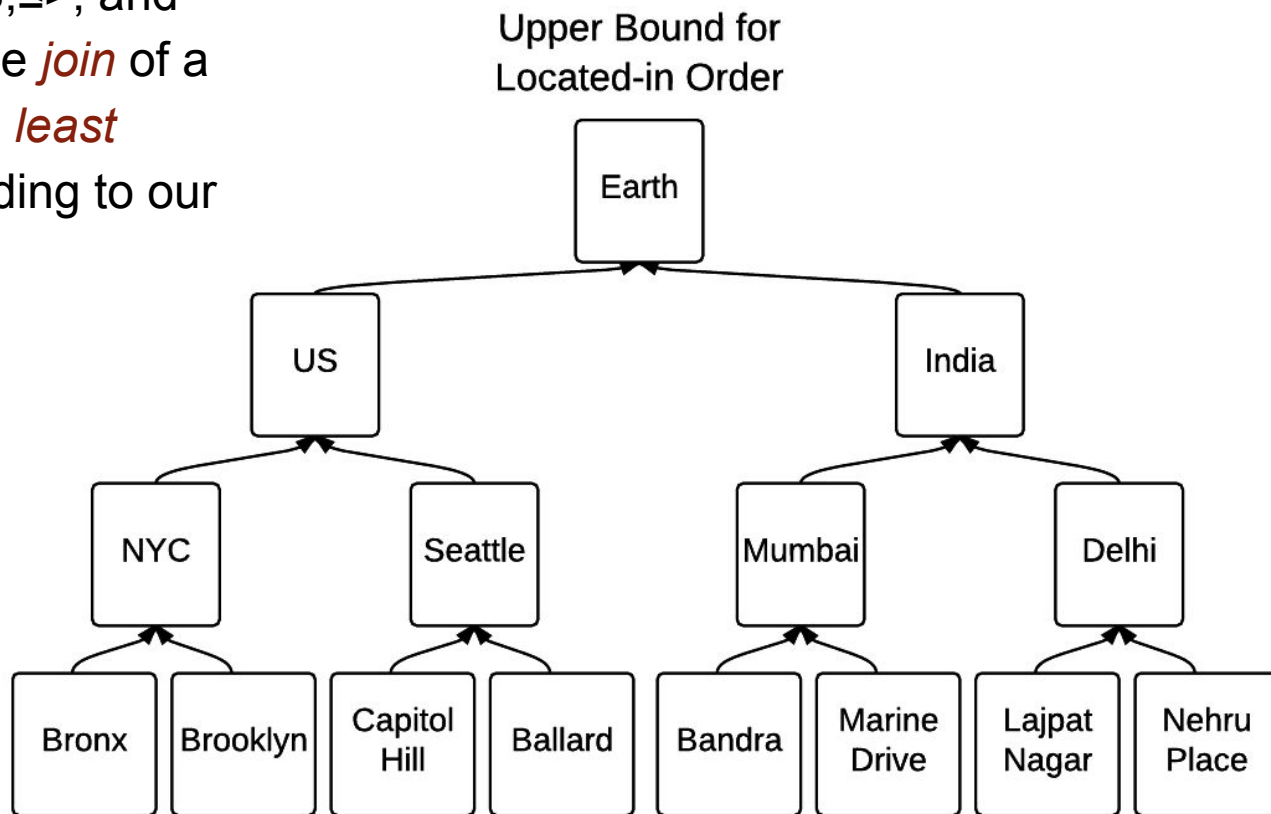
$$(1, 3, 5) \leq (1, 4, 6)$$
$$(2, 7, 2) \leq (8, 7, 3)$$




$$(1, 4, 6) \parallel (2, 7, 2)$$

Join

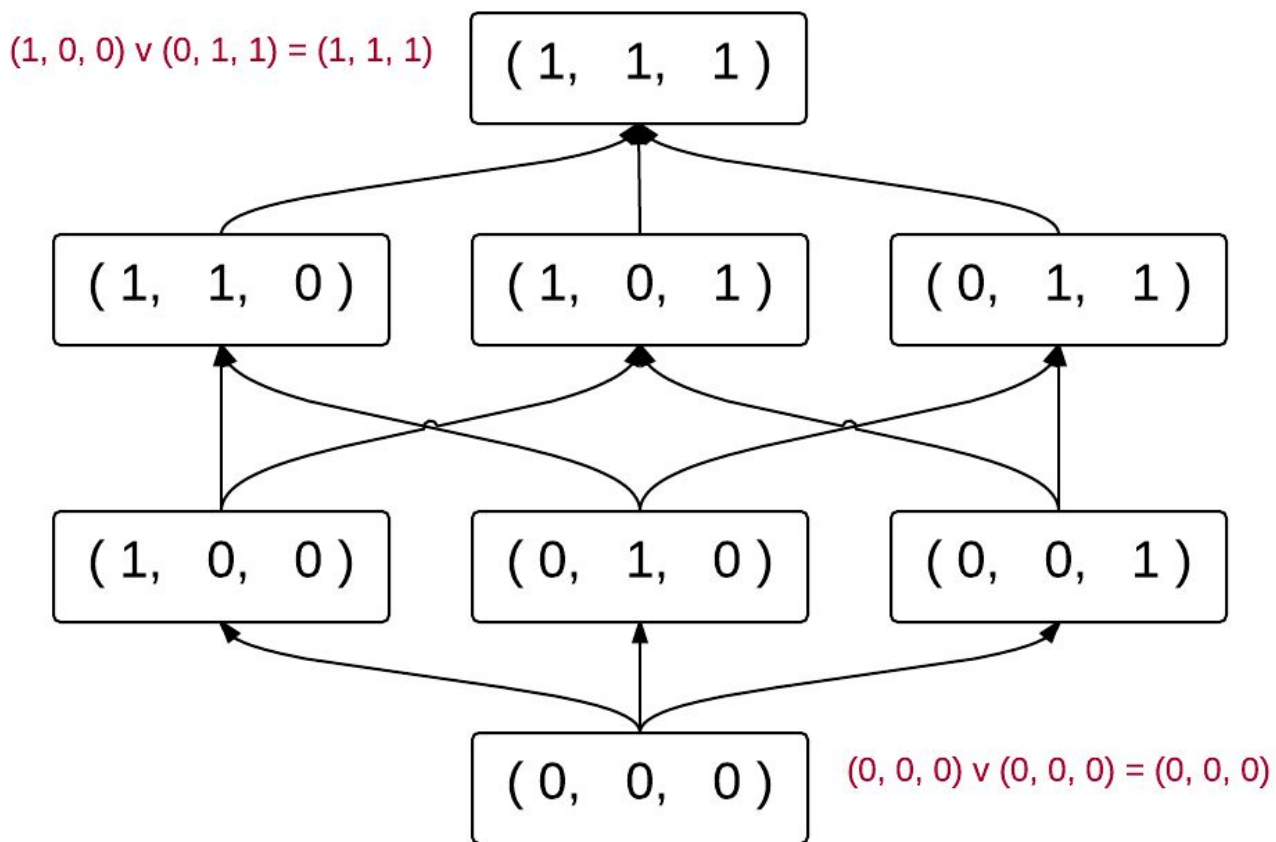
For a set S , an order $\langle S, \leq \rangle$, and two elements $a, b \in S$, the *join* of a and b (written $a \vee b$) is a *least upper bound* of S according to our order $\langle S, \leq \rangle$



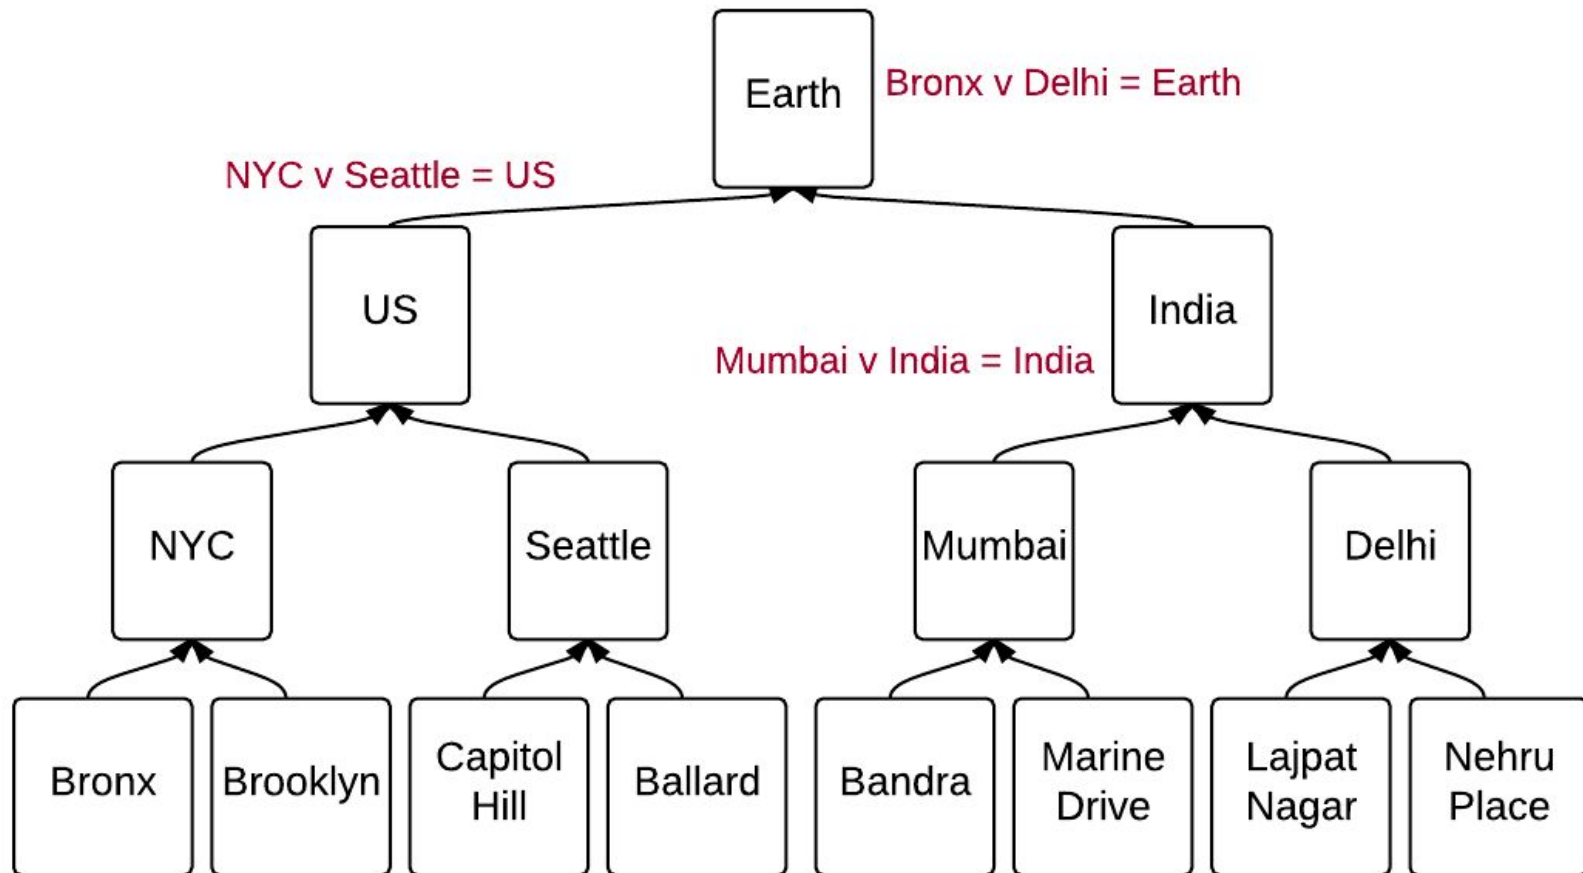
Less Than or Equal To Order



Happened-Before Order




Located-in Order



- A *join semilattice* is an order $\langle S, \leq \rangle$ for which there exists a join $x \vee y$ for any $x, y \in S$


$$\begin{aligned}(0, 0, 0) \vee (0, 0, 1) &= (0, 0, 1) \\ (1, 0, 0) \vee (0, 1, 1) &= (1, 1, 1) \\ (1, 0, 1) \vee (1, 0, 1) &= (1, 0, 1) \\ (0, 1, 0) \vee (0, 0, 1) &= (0, 1, 1)\end{aligned}$$

Happened-Before Order

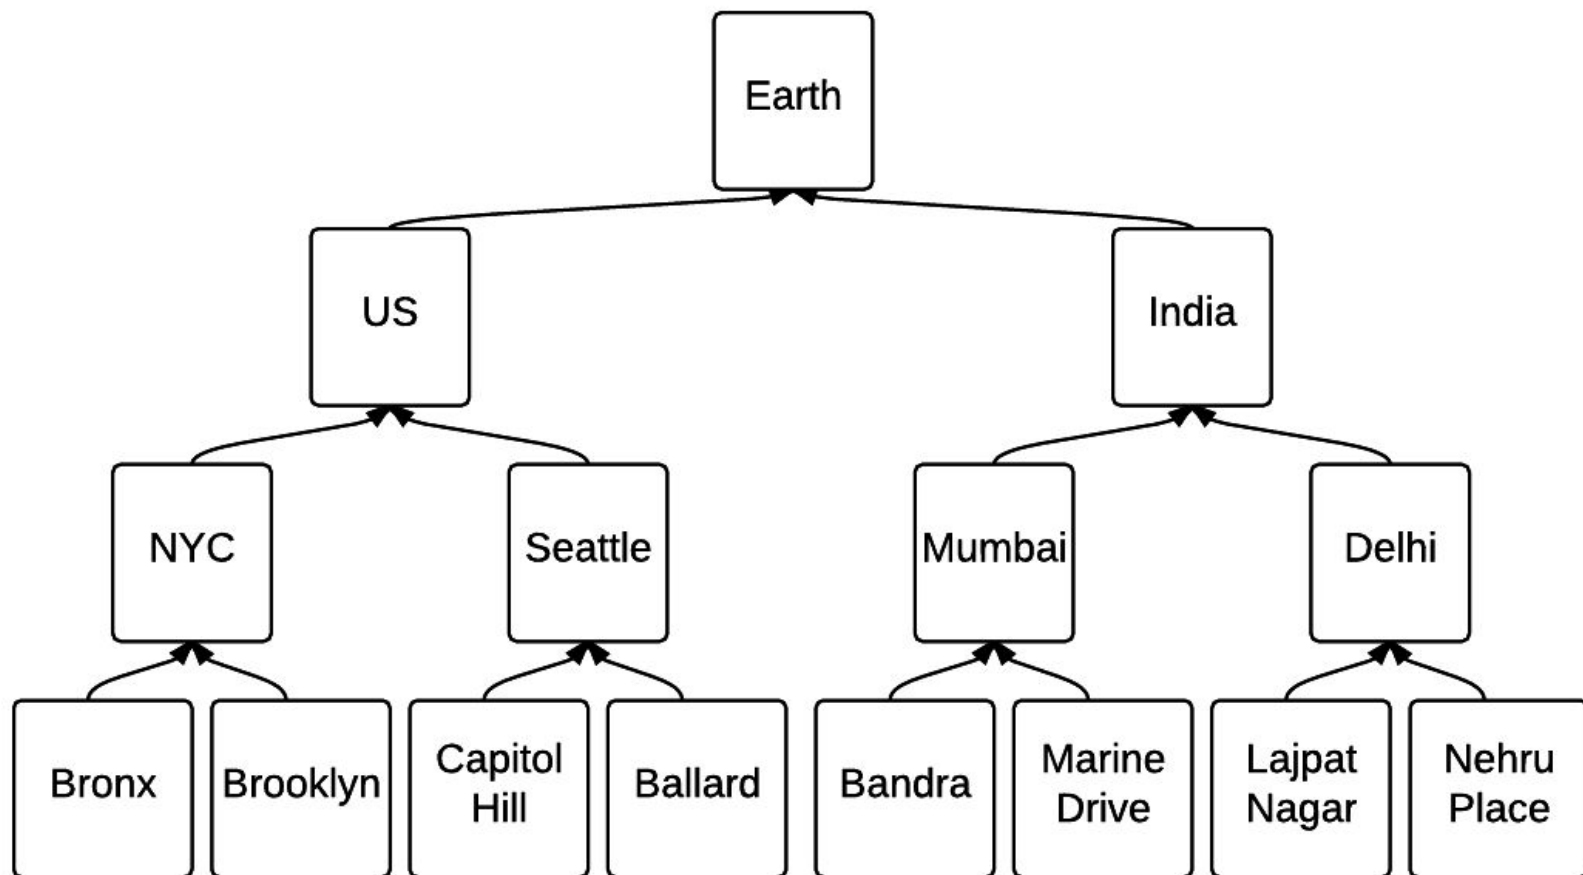


NOT A JOIN SEMI-LATTICE!

$(1, 1, 0) \vee (0, 1, 1)$
does not exist



Located-in Order



Joins obey three laws

- *Commutativity*: $a \vee b = b \vee a$
- *Associativity*: $(a \vee b) \vee c = a \vee (b \vee c)$
- *Idempotence*: $a \vee a = a$

- Joins tend to move *“upwards”*, so do merges of state-based CRDTs tend to converge on the *One True Value*

Convergent CRDTs

- *State* (elements of set)
- *merge()* function

$$\text{merge}(1, 3) = 3$$

$$\text{merge}(9, 5) = 9$$

$$\text{merge}(8, 8) = 8$$

- *merge()* is *max()* here
- Can we use *sum()*?

$$\text{merge}((1, 0, 0), (0, 1, 1)) = (1, 1, 1)$$

$$\text{merge}((0, 0, 0), (2, 0, 2)) = (2, 0, 2)$$

$$\text{merge}((5, 3, 1), (1, 9, 2)) = (5, 9, 2)$$

$$\text{merge}(\text{Seattle}, \text{Mumbai}) = \text{Earth}$$

$$\text{merge}(\text{Bronx}, \text{NYC}) = \text{NYC}$$

$$\text{merge}(\text{Mumbai}, \text{Delhi}) = \text{India}$$

- *System*: set of available state at the moment

[2, 5, 7]

- *Background set*: all integers
- *Value* of the System: upper bound of corresponding semilattice diagram (consistent value)

$\text{Value}([2,5,7])=7$

- The *order of merges* doesn't matter. This is guaranteed by the *associativity* and *commutativity* of joins.
- It doesn't matter how many times we *repeat* a particular merge. This is guaranteed by the *idempotence* of joins.

Why do we care about this ?

Implementing a CvRDT

counter with a simple interface:

- *increment()*: increment the counter
 - *value()*: gets the value of the counter
-
- 3 nodes X, Y, Z
 - Set includes all integers
 - merge() is max()

Imagine the following history:


- Start with 0 on all nodes
- Node 1 increments 3 times
- Node 2 increments 2 times
- Node 3 increments 1 time


What should be the final result ?


Imagine the following history:


- Start with 0 on all nodes
- Node 1 increments 3 times
- Node 2 increments 2 times
- Node 3 increments 1 time


What should be the final result ?













- Weren't we supposed to get **6**?
- Let's use a better approach


- Instead of *integers* use *vector of integers*
- *Value*: sum of all elements in the vectors


Last example becomes:

- X: (3, 0, 0)
- Y: (0, 2, 0)
- Z: (0, 0, 1)











$$\max((3, 0, 0), (0, 2, 1)) = (3, 2, 1)$$


(3, 2, 1)


(0, 2, 1)


(0, 0, 1)

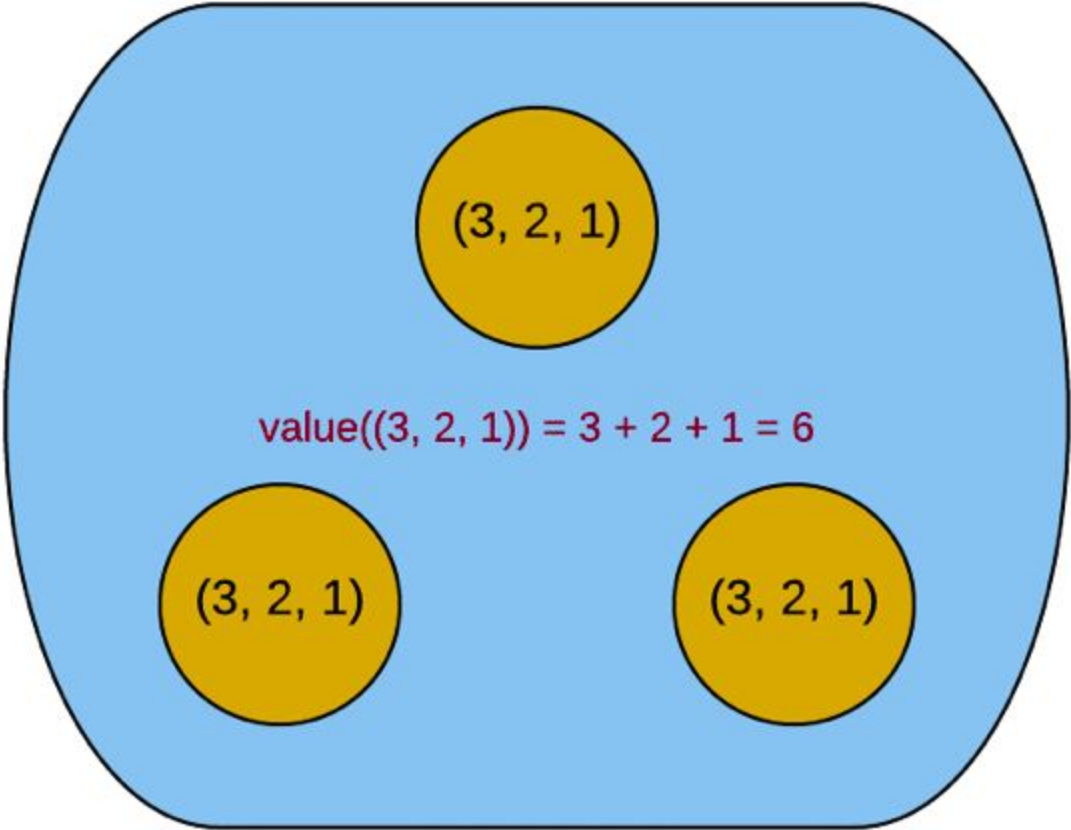













$(3, 2, 1)$

$\text{value}((3, 2, 1)) = 3 + 2 + 1 = 6$

$(3, 2, 1)$

$(3, 2, 1)$

Semi-Lattice for Our System



- Create data-types that follow these *requirements*

References

- <http://jtfmumm.com/blog/2015/11/17/crdt-primer-1-defanging-order-theory/>
- <http://jtfmumm.com/blog/2015/11/24/crdt-primer-2-convergent-crdts/>
- CRDTs: Consistency without concurrency control
<https://arxiv.org/pdf/0907.0929.pdf>
- "CRDTs Illustrated" by Arnout Engelen
<https://www.youtube.com/watch?v=9xFfOhasiOE>
- CRDTs and the Quest for Distributed Consistency by Martin Kleppmann
<https://www.youtube.com/watch?v=B5NULPSiOGw>
- Paxos Simplified <https://www.youtube.com/watch?v=SRsK-ZXTeZ0>
- An extensive list of articles here <https://github.com/ipfs/research-CRDT/>
- https://en.wikipedia.org/wiki/Conflict-free_replicated_data_type